

Introduction to MLOps

Vladislav Goncharenko



MIPT, spring 2026

Vladislav Goncharenko

- Author of machine learning courses and Masters program at MIPT
- Associate professor at MSU, HSE, Harbour.space
- ML researcher (MIPT)
- Team lead of video ranking team at Dzen (yandex.ru)
- Ex-team lead of perception team at self-driving trucks (Evocargo)
- Open source fan



girafe
ai



Дзен

Team

- Demid Gaibov is main seminarist responsible for practical part
- Vladislav Naumov is author of practical sections, seminarist and assistant
- Eugene Astafurov contributed a lot to materials and course structure

Preface

girafe
ai

00

Prerequisites

- UNIX command line basics
 - git basics
- Python programming experience
 - web services is a plus
- Basic machine learning knowledge
 - sklearn
 - pytorch

Disclaimer

- We discuss broad way to do things
- However there are plenty of ways
- We prefer ways suitable for most ML engineers
 - meaning not very deep software engineering skills
- Python is main tool for this course
 - MLOps exists in other stacks too, but for ML it's mostly Python

Your part

1. Knowledge level of each student is different
2. We adjust materials based on your knowledge
3. Your knowledge is estimated by feedback

More you ask - more you get from this course!

I can give you a way but you need to make an effort to follow it

Outline

1. Motivation
2. What is MLOps?
3. Subtasks
 - a. course structure
4. Stacks overview
5. Homeworks outline

Motivation for MLOps


girafe
ai


01

Why MLOps is needed?

- Resources for ML are growing
 - models training (gpu)
 - data mining (cpu, disks)
 - labeling (people)

LANGUAGE MODEL SIZES TO MAR/2023





Why MLOps is needed?

- Resources for ML are growing
 - data mining (cpu)
 - labeling (people)
 - models training (gpu)
- Results reproduction
 - not only for industry, but for academia as well

Links: [one](#), [two](#), [three](#), [four](#)

Why Most Published Research Findings Are False

🌐 1 language ▼

Article [Talk](#)

[Read](#) [Edit](#) [View history](#) [Tools](#) ▼

From Wikipedia, the free encyclopedia

COMMUNICATIONS OF THE ACM

[HOME](#)

[CURRENT ISSUE](#)

[NEWS](#)

[BLOGS](#)

[OPINION](#)

[RESEARCH](#)

[PRACTICE](#)

[CAREERS](#)

[ARCHIVE](#)

[VIDEOS](#)

Search

"Why Most Published Research Findings Are False" by Ioannidis et al. from the [Stanford School of Medicine](#) [metascience](#).

In the paper, Ioannidis argues that many published research findings contain results that cannot be replicated. He determines whether scientific findings are based on formalized calculations ("P values") or on untested, unproven assumptions about the world, which indicates that most published research findings are false.

[Home](#) / [Magazine Archive](#) / [August 2020 \(Vol. 63, No. 8\)](#) / [Threats of a Replication Crisis in Empirical Computer Science](#) / [Abstract](#)

REVIEW ARTICLES

Threats of a Replication Crisis in Empirical Computer Science

By Andy Cockburn, Pierre Dragicevic, Lonni Besançon, Carl Gutwin

Communications of the ACM, August 2020, Vol. 63, No. 8, Pages 70-79


Why MLOps is needed?

- Resources for ML are growing
 - data mining (cpu)
 - labeling (people)
 - models training (gpu)
- Results reproduction
 - not only for industry, but for academia as well
- Models delivery
 - reduce time to market
 - reduce routine
 - testing and reliability
- Competence decomposition
 - deeper division of labor

Why do ML specialists need MLOps?

The Virgin Statistics

Entire skillset falls to pieces the moment they work with an uncommon distribution



Obsessed with plotting. Claims it helps people understand, but it only confuses them more.



Convolved, ineffective methods get rebuffed at every work meeting.

"I k...know the results seem weird, b-but the p-value is considered significant..."

Decisions:	The Null Hypothesis is True	False
Accept H_0	(1- α) Confidence	β
Reject H_0	α	(1- β) Power of Test

Checks his results.

Wears suits and cologne to seem more professional

Obsessed with theory. Can't do anything without stealing the methods of smarter, dead men.



Secretly resentful, wishes his field was more popular.

Constantly begging superiors for more data.

Knows *maybe* three methods, all requiring impossible assumptions.

Simple Linear Regression Model

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

Assumptions: Random, Independent, Homoscedastic, Normality


The Chad Machine-Learning

Has literally never heard the word "distribution" in his life



Regularizes his data until it gives the result he wants.

Fastest growing and most in demand discipline in his country.





Innovates daily, advancing his field while solving real problems.

Doesn't try to explain methods, simply gives results, and is respected for it.

Wears sweats and muscle shirts to work, no one can stop him because he's too valuable.

Has so much data, often has to throw it away.

Let's his subordinates waste their days making his plots, would be a waste of his talent.



Refuses to use the same method more than once, each problem is unique and requires its own touch.

Implements methods into production without testing.

Why do ML specialists need MLOps?

because no one will do it for you!!!

For deeper dive I suggest reviewing
system design primer

<https://github.com/donnemartin/system-design-primer>

- cookbook for modern web based systems

Greg Brockman, OpenAI CTO

“... we need more people with strong software skills who can deliver these systems. This is because we are building AI models out of unprecedented amounts of compute...”

“Harnessing this compute requires deep software skills and the right kind of machine learning knowledge. We need to coordinate lots of computers, build software frameworks”

<https://blog.gregbrockman.com/its-time-to-become-an-ml-engineer>



What is MLOps?

girafe
ai

02


DevOps definition

DevOps is a methodology in the software development and IT industry. Used as a set of practices and tools.

DevOps integrates and automates the work of software development (Dev) and IT operations (Ops) as a means for improving and shortening the systems development life cycle

[Well-known common knowledge site](#)

[Great talk on DevOps practices](#)





[DevOps course at MIPT \(ru\)](#)

MLOps definition

MLOps is a paradigm that aims to deploy and maintain machine learning models in production reliably and efficiently.


MLOps seeks to increase automation and improve the quality of production models, while also focusing on business and regulatory requirements

[Well-known common knowledge site](#)



MLOps = DataML + DevOps

ml4devs.com/mlops-lifecycle 



© 2022 Satish Chandra Gupta




CC BY-NC-ND 4.0 International License
creativecommons.org/licenses/by-nc-nd/4.0/

scgupta.me 
twitter.com/scgupta 
linkedin.com/in/scgupta 



Relations to other spheres

- Management
- System Design
- Software development
- Data bases
 - SQL
- Analytics
- Hardware
 - Cluster design
- Machine Learning



Processes and management


girafe
ai

03

CRISP DM

In the field of Data Science is used. **Cross-industry standard process for data mining** ([CRISP DM](#)) is a process standard for projects using predictive analytics, proposed in 1999. [It was used by half of the companies](#) engaged in Data Mining at the beginning of the century.


[Read more](#)



Poll Years	2002	2004	2007	2014
CRISP-DM	51%	42%	42%	43%
SEMMA	12%	10%	13%	8.5%
KDD process			7%	7.5%
My organization's	7%	6%	5%	3.5%
My own	23%	28%	19%	27.5%
Other (incl. domain specific)	4%	6%	9% (5%)	10% (2%)
None	4%	7%	5%	0%

Team Data Science Process

Data Science Lifecycle




Modern standard from Microsoft

<https://learn.microsoft.com/en-us/azure/architecture/data-science-process/overview>

It is distinguished by the allocation of roles and a more detailed study of the stages of the cycle


TDSP: team roles



Other approaches


1. [SEMMA](#)
[read more](#)
2. [KDD Process](#)
3. [Waterfall](#)
4. etc...

In fact these are topics for IT management course



MLOps gradual introduction steps

It's very important to go by MLOps road leisurely because change of processes is always painful for people accustomed to current instruments and solutions



MLOps as a requirement for ML

<https://proceedings.neurips.cc/paper/2015/file/86df7dcfd896fcdf2674f757a2463eba-Paper.pdf>

Hidden Technical Debt in Machine Learning Systems

D. Sculley, Gary Holt, Daniel Golovin, Eugene Davydov, Todd Phillips
{dsculley, gholt, dgg, edavydov, toddphillips}@google.com
Google, Inc.




Figure 1: Only a small fraction of real-world ML systems is composed of the ML code, as shown by the small black box in the middle. The required surrounding infrastructure is vast and complex.

Further reading

- <https://ml-ops.org/content/mlops-principles>
- <https://www.ml4devs.com/en/articles/mlops-machine-learning-life-cycle/>
- <https://ashutoshtripathi.com/2021/08/18/mlops-a-complete-guide-to-machine-learning-operations-mlops-vs-devops/>
- <https://netraneupane.medium.com/introduction-to-mlops-for-beginner-part-1-37207ea3004a>
- <https://neptune.ai/blog/how-to-learn-mlops>

MLOps subtasks

Course structure

girafe
ai

03

MLOps subtasks

- Contemporary solutions overview
- Software engineering
 - Python infrastructure
- Virtualization
 - Docker
- Code storage and management
 - Git
 - CI & CD
 - code quality
 - testing (auto tests, test stands)
 - documentation generation
- Data storage and management
 - Data Acquisition
 - Data Engineering
 - AB tests

MLOps subtasks

- Contemporary solutions overview
- Software engineering
- Virtualization
- Code storage and management
- Data storage and management
- Labeling and crowdsourcing
- Experiments tracking
 - production monitoring
- Computational resources
 - GPU
- Model storages and deployment preparation
- Efficient inference
- Embedding storage optimization and vector databases
- LLMOps: vLLM, ollama and friends



MLOps stacks overview

girafe
ai

04

MLOps at big tech

Big companies has their stack to organize models training and publishing pipelines

- Bigtech
 - [Amazon SeageMaker](#), [example](#)
 - [Google Cloud](#)
 - [Nvidia](#)
 - [Microsoft Azure](#)
 - Yandex: [YTSaurus](#) + Nirvana + [Toloka](#)
- Independent vendors
 - [Databricks](#)
 - [Seldon](#)

Generally variance of instruments is huge over different companies (from only bare metal ssh machines to very solid custom solutions)


Closed circle solutions

Pros:

- high level of integration
- reliability

Cons:

- high price
- binding to vendor
 - vendorlock



Not only big tech


[Youtube video](#)

**DATA+AI
SUMMIT**
BY  databricks


MLOps at Gucci: from zero to hero

An overview on implementing an MLOps solution from scratch


Databricks
2023



Paid digital media landscape is vast




* террористы



Data driven decision

How to allocate budget on media channels to maximize ROI?


Budget scenarios



GROAS 2.0




Net sales forecasts



Best scenario is chosen for future allocation

Get faster to production

GROAS 1.0
6 months (*)



GROAS 2.0
6 weeks (*)



MLOps at academia

It just doesn't exist (unfortunately)

You publish or perish:
no time for bullshit



What about open source stack?

Will discuss it during all the course.

Main problem of FOSS instruments - **lack of orchestration**.

Each instrument primarily aims to solve one task but then start to add some default solutions for adjacent problems which are not that great.

All in all no instrument solve every problem.

Finally you need to **write glue code yourself**.

Homeworks

girafe
ai

05

Homeworks

1. Formulate problem to apply MLOps stack
2. Write baseline code
3. Apply parts of pipeline from classes
4. Peer review
5. Final review

Also we will have simple tests for each lecture.

One person does one project, no grouping in this course.


Bonus task: create pull request with some mlops instrument to any significant open source project e.g. [python-dotenv](#) lacks contemporary dependency management tool

Resources on MLOps

- <https://github.com/DataTalksClub/mlops-zoomcamp>
- <https://github.com/mryab/efficient-dl-systems/tree/main>
-

Revise

- Motivation
- What is MLOps?
- Subtasks
 - and course structure
- Stacks overview
- Homeworks outline



Thanks for attention!

Questions?

